


# Document Similarity

Natural Language Processing



# Agenda

---

- What is Document Similarity
- Methods to measure Document Similarity
- Cosine Similarity Method

# Goal

---

- Given a set of documents and search term(s)/query we need to retrieve relevant documents that are **similar to the search query**.

# Statistical Retrieval

---

- Retrieval based on **similarity** between query and documents.
- Output documents are ranked according to similarity to query.
- Similarity based on occurrence frequencies of keywords in query and document.

# The Vector-Space Model

---

- Assume  $t$  distinct terms remain after preprocessing; call them index terms or the vocabulary.
- These terms form a vector space.
  - Dimensionality =  $t = |\text{vocabulary}|$
- Each term,  $i$ , in a document or query,  $j$ , is given a real-valued weight,  $w_{ij}$ .
- Both documents and queries are expressed as  $t$ -dimensional vectors:
  - $d_j = (w_{1j}, w_{2j}, \dots, w_{tj})$

# Issues for Vector Space Model

---

- How to determine important words in a document?
  - Word sense?
  - Word n-grams (and phrases, idioms,...)
- How to determine the degree of importance of a term within a document and within the entire collection?
- How to determine the degree of similarity between a document and the query?
- In the case of the web, what is the collection and what are the effects of links, formatting information, etc.?


# Graphic Representation

Example:

$$D_1 = 2T_1 + 3T_2 + 5T_3$$

$$D_2 = 3T_1 + 7T_2 + T_3$$

$$Q = 0T_1 + 0T_2 + 2T_3$$



- Is  $D_1$  or  $D_2$  more similar to  $Q$ ?
- How to measure the degree of similarity? Distance? Angle? Projection?

# Document Collection

A collection of  $n$  documents can be represented in the vector space model by a term-document matrix.

An entry in the matrix corresponds to the “weight” of a term in the document; zero means the term has no significance in the document or it simply doesn't exist in the document.

$$\begin{pmatrix} & T_1 & T_2 & \dots & T_t \\ D_1 & w_{11} & w_{21} & \dots & w_{t1} \\ D_2 & w_{12} & w_{22} & \dots & w_{t2} \\ \vdots & \vdots & \vdots & & \vdots \\ \vdots & \vdots & \vdots & & \vdots \\ D_n & w_{1n} & w_{2n} & \dots & w_{tn} \end{pmatrix}$$



# Term Weights: Term Frequency

---

- More frequent terms in a document are more important, i.e. more indicative of the topic.

$f_{ij}$  = frequency of term  $i$  in document  $j$

- May want to normalize *term frequency* ( $tf$ ) by dividing by the frequency of the most common term in the document:

$$tf_{ij} = f_{ij} / \max_i \{f_{ij}\}$$

# Term Weights: Inverse Document Frequency

---

- Terms that appear in many *different* documents are *less* indicative of overall topic.

$df_i$  = document frequency of term  $i$   
= number of documents containing term  $i$

$idf_i$  = inverse document frequency of term  $i$ ,  
=  $\log_2 (N / df_i)$   
( $N$ : total number of documents)

- An indication of a term's *discrimination* power.
- Log used to dampen the effect relative to  $tf$ .

# TF-IDF Weighting

- A typical combined term importance indicator is *tf-idf weighting*:

$$w_{ij} = tf_{ij} idf_i = tf_{ij} \log_2 (N/ df_i)$$

- A term occurring frequently in the document but rarely in the rest of the collection is given high weight.
- Many other ways of determining term weights have been proposed.
- Experimentally, *tf-idf* has been found to work well.

# Computing TF-IDF -- An Example

---

- Given a document containing terms with given frequencies:
- $A(3), B(2), C(1)$
- Assume collection contains 10,000 documents and
- document frequencies of these terms are:
- $A(50), B(1300), C(250)$
- Then:
- A:  $tf = 3/3$ ;  $idf = \log_2(10000/50) = 7.6$ ;  $tf-idf = 7.6$
- B:  $tf = 2/3$ ;  $idf = \log_2(10000/1300) = 2.9$ ;  $tf-idf = 2.0$
- C:  $tf = 1/3$ ;  $idf = \log_2(10000/250) = 5.3$ ;  $tf-idf = 1.8$

# Query Vector

---

- Query vector is typically treated as a document and also tf-idf weighted.
- Alternative is for the user to supply weights for the given query terms.

# Similarity Measure

---


- A **similarity measure** is a function that computes the **degree of similarity** between two vectors.
- Using a similarity measure between the query and each document:
- It is possible to rank the retrieved documents in the order of presumed relevance.
- It is possible to enforce a certain threshold so that the size of the retrieved set can be controlled.

# Cosine Similarity Measure

Cosine similarity measures the cosine of the angle between two vectors.

Inner product normalized by the vector lengths.

$$\text{CosSim}(\mathbf{d}_j, \mathbf{q}) = \frac{\vec{d}_j \cdot \vec{q}}{|\vec{d}_j| \cdot |\vec{q}|} = \frac{\sum_{i=1}^t (w_{ij} \cdot w_{iq})}{\sqrt{\sum_{i=1}^t w_{ij}^2 \cdot \sum_{i=1}^t w_{iq}^2}}$$



$$\begin{aligned} D_1 &= 2T_1 + 3T_2 + 5T_3 & \text{CosSim}(D_1, Q) &= 10 / \sqrt{(4+9+25)(0+0+4)} = 0.81 \\ D_2 &= 3T_1 + 7T_2 + 1T_3 & \text{CosSim}(D_2, Q) &= 2 / \sqrt{(9+49+1)(0+0+4)} = 0.13 \\ Q &= 0T_1 + 0T_2 + 2T_3 \end{aligned}$$

$D_1$  is 6 times better than  $D_2$  using cosine similarity

# Naïve Implementation

- Convert all documents in collection D to *tf-idf* weighted vectors,  $\mathbf{d}_j$ , for keyword vocabulary V.
- Convert query to a *tf-idf*-weighted vector  $\mathbf{q}$ .
- For each  $\mathbf{d}_j$  in D do
  - Compute score  $s_j = \text{cosSim}(\mathbf{d}_j, \mathbf{q})$
- Sort documents by decreasing score.
- Present top ranked documents to the user.
- Time complexity:  $O(|V| \cdot |D|)$  Bad for large V & D !
- $|V| = 10,000$ ;  $|D| = 100,000$ ;  $|V| \cdot |D| = 1,000,000,000$



# Comments on Vector Space Models

---

- Simple, mathematically based approach.
- Considers both local (tf) and global (idf) word occurrence frequencies.
- Provides partial matching and ranked results.
- Tends to work quite well in practice despite obvious weaknesses.
- Allows efficient implementation for large document collections.

# Problems with Vector Space Model

---

- Missing semantic information
  - word sense
- Missing syntactic information
  - phrase structure, word order, proximity information
- Assumption of term independence
  - ignores synonymy
- Given a two-term query “A B”
  - may prefer a document containing A frequently but not B, over a document that contains both A and B, but both less frequently.